

ASSIGNMENT 4

Problem Statement:

We have a dataset of sales of different TV sets across different locations.

Records look like:

Samsung|Optima|14|Madhya Pradesh|132401|14200

The fields are arranged like:

Company Name|Product Name|Size in inches|State|Pin Code|Price

There are some invalid records which contain 'NA' in either Company Name or Product Name.

TASK 1:

Write a Map Reduce program to filter out the invalid records. Map only job will fit for this context.

SOLUTION:

For this task I created a method that checks if a record is invalid. To achieve this, we split every record using “\” character as rule, into an array. Then we check every part of the array and we validate if that part contains “NA”, then we return a Boolean variable:

```
private boolean recordIsValid(Text record){
    String[] lineArray = record.toString().split("\\|");
    boolean isValid = true;
    for(int i=0; i<lineArray.length; i++){
        if(lineArray[i].equals("NA")){
            isValid = false;
        }
    }
    return isValid;
}
```

The mapper uses this method to filter out invalid records:

```
public void map(LongWritable key, Text value, Context context)
throws IOException, InterruptedException{
    if(recordIsValid(value)==false){
        Text record = new Text();
        record = value;
        context.write(key,record);
    }
}
```

Then we need to export the .jar file from Eclipse and in the command shell we type:

hadoop jar /home/acadgild/workspace/TvSales.jar television.txt /skip-bad-records

- **jar** is the command to execute the mapping and reducing
- **/home/acadgild/workspace/TvSales.jar** location on file system of the jar exported before
- **television.txt** is the file location on HDFS containing the records
- **/skip-bad-records** is the destination directory on HDFS after reducing the records. Since there is no reducer for this task, Hadoop uses identity Reducer .

To check the results obtained on HDFS we can type:

hadoop fs -cat /filter-bad-records/part-r-00000

```
File Edit View Search Terminal Help
File Input Format Counters
  Bytes Read=733
File Output Format Counters
  Bytes Written=706
[acadgild@localhost workspace]$ hadoop fs -cat /skip-bad-records/part-r-00000
18/07/24 17:34:55 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
0      Samsung|Optima|14|Madhya Pradesh|132401|14200
47     Onida|Lucid|18|Uttar Pradesh|232401|16200
90     Akai|Decent|16|Kerala|922401|12200
126    Lava|Attention|20|Assam|454601|24200
164    Zen|Super|14|Maharashtra|619082|9200
202    Samsung|Optima|14|Madhya Pradesh|132401|14200
249    Onida|Lucid|18|Uttar Pradesh|232401|16200
292    Onida|Decent|14|Uttar Pradesh|232401|16200
369    Lava|Attention|20|Assam|454601|24200
407    Zen|Super|14|Maharashtra|619082|9200
445    Samsung|Optima|14|Madhya Pradesh|132401|14200
532    Samsung|Decent|16|Kerala|922401|12200
571    Lava|Attention|20|Assam|454601|24200
609    Samsung|Super|14|Maharashtra|619082|9200
651    Samsung|Super|14|Maharashtra|619082|9200
693    Samsung|Super|14|Maharashtra|619082|9200
[acadgild@localhost workspace]$
```

The complete code is in **filter-invalid-records-program.txt** and the filtered records in **filtered-records.txt** obtained from HDFS using **hadoop fs -copyToLocal** command.

TASK 2:

Write a Map Reduce program to calculate the total units sold for each Company.

SOLUTION:

To calculate the total units sold for each, we need to do “the sort and shuffle” in the mapper class method. For this, we split the records, and we select the “key” as company and the value as “1” as integer, so for every record there’s a 1 value. By the way, we use the method to filter out bad records.

```
public void map(LongWritable key, Text value, Context context)
throws IOException, InterruptedException{

    if(recordIsValid(value)==false){
        Text company = new Text();
        IntWritable unit = new IntWritable();
        company = new Text(value.toString().split("\\|")[0]);
        unit = new IntWritable(1);
        context.write(company, unit );
    }
}
```

Now we need to reduce the data in the Reducer class method. The objective is to sum every value for each key, with this; we will get the total units sold by each company. To achieve this we write the following code:

```
private IntWritable result = new IntWritable();

public void reduce (Text key, Iterable<IntWritable> values, Context context)
throws IOException, InterruptedException{
    int sum = 0;
    for(IntWritable val: values){
        sum += val.get();
    }
    result.set(sum);
    context.write(key, result);
}
```

Then we need to export the .jar file from Eclipse and in the command shell we type:

```
hadoop jar /home/acadgild/workspace/TvSales.jar television.txt  
/total-units-sold
```

- **jar** is the command to execute the mapping and reducing
- **/home/acadgild/workspace/TvSales.jar** location on file system of the jar exported before.
- **television.txt** is the file location on HDFS containing the records
- **/total-units-sold** is the destination directory on HDFS after reducing the records.

To check the results obtained on HDFS we can type :

hadoop fs -cat /total-units-sold/part-r-00000

```

Physical memory (bytes) snapshot=268697600
Virtual memory (bytes) snapshot=4113760256
Total committed heap usage (bytes)=137498624
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
Bytes Read=733
File Output Format Counters
Bytes Written=38
[acadgild@localhost ~]$ hadoop fs -cat /total-units-sold/part-r-00000
18/07/26 01:37:34 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Akai      1
Lava      3
Onida     3
Samsung   7
Zen       2

```

The complete code is in **total-units-sold-program.txt** and the results are in **total-units-sold.txt** obtained from HDFS using **hadoop fs -copyToLocal** command.

In this case, we will use the same reduce method than before, but we need to change the mapper method.

```

public void map(LongWritable key, Text value, Context context)
throws IOException, InterruptedException{

    if(recordIsValid(value)==false & value.toString().split("\\|")[0].equals("onida")){
        Text state = new Text();
        IntWritable unit = new IntWritable();
        state = new Text(value.toString().split("\\|")[3]);
        unit = new IntWritable(1);
        context.write(state, unit );
    }
}

```

So besides filtering out invalid records, we filter records by company. In this case, we will only accept records with “Onida” as a company. As a key, we will assign the state, so we can sum every record by state, and the value the same as before.

Then we need to export the .jar file from Eclipse and in the command shell we type:

```
hadoop jar /home/acadgild/workspace/TvSales.jar television.txt  
/total-units-onida
```

- **jar** is the command to execute the mapping and reducing
- **/home/acadgild/workspace/TvSales.jar** location on file system of the jar exported before
- **television.txt** is the file location on HDFS containing the records
- **/total-units-onida** is the destination directory on HDFS after reducing the records.

To check the results obtained on HDFS we can type

```
hadoop fs -cat /total-units-onida/part-r-00000
```

```
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=464
CPU time spent (ms)=2590
Physical memory (bytes) snapshot=267829248
Virtual memory (bytes) snapshot=4113756160
Total committed heap usage (bytes)=137498624

Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=733
File Output Format Counters
  Bytes Written=16
[acadgild@localhost ~]$ hadoop fs -cat /total-units-onida/part-r-00000
18/07/26 08:51:46 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Uttar Pradesh 3
[acadgild@localhost ~]$ █
```

We only get one state, because the other records are invalid.

The complete code is in **total-units-onida-program.txt** and the results are in **total-units-onida.txt** obtained from HDFS using **hadoop fs -copyToLocal** command